

Exploratory data analysis

- missingness audit (stacked bar charts and correlation heatmap)
- z-score sanity check (histograms)
- demographic breakdown (bar charts and count plots)
- severity distributions (stacked bar charts)
- academic performance vs. mental health (violin plots)
- timeline of burnout (line plots)
- inter-condition correlations (correlation matrix heatmap)
- z-score density by source (overlapping density plots)
- departmental pressure (horizontal boxplots)
- multivariate density comparison of conditions (corner pairplot)

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from pathlib import Path

from utils import style_pct
```

global constants

```
# assume notebook is run from proj_root/mhp/
ROOT = Path.cwd().parent
DATA = ROOT / "mhp" / "data"
PLOTS = ROOT / "mhp" / "plots"

Z_SCORE_COLS = ["stress_z_score", "anxiety_z_score", "depression_z_score"]
SEVERITY_COLS = ["stress_level", "anxiety_level", "depression_level"]
DEMOGRAPHIC_COLS = ["age", "gender", "academic_year", "cgpa"]

PLACEHOLDER_LABELS = ["other", "unknown", "not_recorded"]
SEVERITY_ORDER = ["none", "minimal", "mild", "moderate", "moderately_severe", "severe"]

CLEAN_LABELS = {
    "age": "Age",
    "gender": "Gender",
    "academic_year": "Academic Year",
    "cgpa": "CGPA",
    "stress_z_score": "Stress (Z-Score)",
    "anxiety_z_score": "Anxiety (Z-Score)",
    "depression_z_score": "Depression (Z-Score)",
    "stress_level": "Stress (Level)",
    "anxiety_level": "Anxiety (Level)",
    "depression_level": "Depression (Level)",
    "dataset1": "Primary Dataset",
    "dataset2": "Secondary Dataset",
}
```

```
# color scheme
COLOR_PRIMARY = "steelblue"
COLOR_SECONDARY = "sandybrown"
COLOR_NEUTRAL = "lightgray"
CMAP_DIVERGING = "vlag"
CMAP_SEQUENTIAL = "mako"

MISSINGNESS_PALETTE = {
    "NA": "indianred",
    "other": "steelblue",
    "unknown": "sandybrown",
    "not_recorded": "gray",
}
```

```
mhp_dataset = pd.read_parquet(DATA / "mhp_dataset.parquet")
```

missingness audit

```
n = len(mhp_dataset)
na_counts = mhp_dataset.isna().sum()
na_pct = (na_counts / n) * 100

if na_counts.sum() > 0:
    print("Columns with missing values:")
    na_df = pd.DataFrame(
        {"n_missing": na_counts, "pct_missing": (na_counts / n) * 100}
    )
    na_df = na_df[na_df["n_missing"] > 0].sort_values("n_missing")
    display(style_pct(na_df))

sentinel_rows = []
for col in mhp_dataset.columns:
    for label in PLACEHOLDER_LABELS:
        count = mhp_dataset[col].isin([label]).sum()
        if count > 0:
            sentinel_rows.append(
                {"column": col, "level": label, "n": count, "pct": (count / n) * 100}
            )

if sentinel_rows:
    print("Sentinel categories summary")
    sentinel_tbl = pd.DataFrame(sentinel_rows).sort_values(
        ["pct", "level"], ascending=False, ignore_index=True
    )
    display(style_pct(sentinel_tbl))

stack_components = pd.DataFrame(index=mhp_dataset.columns)
stack_components["NA"] = na_pct
for label in PLACEHOLDER_LABELS:
    stack_components[label] = (mhp_dataset.isin([label]).sum() / n) * 100

cols_order = stack_components.sum(axis=1).sort_values(ascending=True).index
```

```

stack_plot = stack_components.loc[cols_order]

fig, (ax_bar, ax_heat) = plt.subplots(1, 2, figsize=(16, 6), constrained_layout=True)

stack_plot.plot(kind="barh", stacked=True, color=MISSINGNESS_PALETTE, ax=ax_bar)

ax_bar.set_xlabel("% of Dataset")
ax_bar.set_title("Missingness and Explicit Sentinels by Column")
ax_bar.legend(loc="lower right")

combined_mask = mhp_dataset.isna() | mhp_dataset.isin(PLACEHOLDER_LABELS)
varied_cols = [col for col in combined_mask.columns if combined_mask[col].var() > 0]
correlation_matrix = combined_mask[varied_cols].corr()

sns.heatmap(
    correlation_matrix,
    ax=ax_heat,
    cmap=CMAP_DIVERGING,
    center=0,
    vmin=-1,
    vmax=1,
    square=True,
    annot=True,
    fmt=".2f",
    cbar_kws={"shrink": 0.8},
)
ax_heat.set_title("Correlation of Missing/Sentinel Co-occurrence")

plt.show()


```

Columns with missing values:

	n_missing	pct_missing
anxiety_level	8	0.34%
anxiety_z_score	8	0.34%

Sentinel categories summary

	column	level	n	pct
0	department	not_recorded	351	14.75%
1	cgpa	other	181	7.61%
2	academic_year	other	147	6.18%
3	department	other	79	3.32%
4	gender	other	10	0.42%
5	academic_year	unknown	5	0.21%



```
fig.savefig(PLOTS / "missingness_audit.svg", format="svg", bbox_inches="tight")
```

z-score sanity check

```
z_summary = (
    mhp_dataset[Z_SCORE_COLS].agg(["count", "mean", "std", "min", "max"]).T
)
print("Z-score summary (expect mean ~0 and std ~1):")
display(
    z_summary.style
    .format("{:.0f}", subset=["count"])
    .format("{:.3f}", subset=["mean", "std", "min", "max"])
)

fig, axes = plt.subplots(
    1, len(Z_SCORE_COLS), figsize=(16, 5), constrained_layout=True
)

for ax, col in zip(axes, Z_SCORE_COLS):
    unique_vals = mhp_dataset[col].nunique()


    sns.histplot(
        mhp_dataset[col].dropna(), bins=unique_vals, kde=True, ax=ax, color=COLOR_PRI
    )

    ax.axvline(0, color="black", linestyle="--", linewidth=1)
    ax.set_title(CLEAN_LABELS[col])
    ax.set_xlabel("Z-score")
    ax.set_ylabel("Count")

fig.suptitle("Z-score Distributions: Stress, Anxiety, Depression")
plt.show()
```

Z-score summary (expect mean ~0 and std ~1):

	count	mean	std	min	max
stress_z_score	2379	-0.000	1.000	-3.405	3.197
anxiety_z_score	2371	0.000	1.000	-2.247	3.061
depression_z_score	2379	-0.000	1.000	-2.164	3.195



```
fig.savefig(PLOTS / "z_score_distributions.svg", format="svg", bbox_inches="tight")
```

demographic breakdown

```
print("Demographic Frequency Tables:")
for col in DEMOGRAPHIC_COLS:
    freq = mhp_dataset[col].value_counts().rename_axis(col).reset_index(name="count")
    freq["pct"] = (freq["count"] / len(mhp_dataset)) * 100
    display(style_pct(freq))

fig, axes = plt.subplots(2, 2, figsize=(14, 8), constrained_layout=True)
axes = axes.flatten()

for ax, col in zip(axes, DEMOGRAPHIC_COLS):
    is_categorical = isinstance(mhp_dataset[col].dtype, pd.CategoricalDtype)
    plot_order = None if is_categorical else mhp_dataset[col].value_counts().index

    sns.countplot(data=mhp_dataset, y=col, order=plot_order, ax=ax, color=COLOR_PRIMARY)
    ax.set_title(CLEAN_LABELS[col])
    ax.set_xlabel("Count")
    ax.set_ylabel("")

fig.suptitle("Demographic Breakdown")
plt.show()
```


Demographic Frequency Tables:

	age	count	pct
0	18-22	1514	63.64%
1	23-26	830	34.89%
2	27-30	28	1.18%
3	<18	4	0.17%
4	>30	3	0.13%

	gender	count	pct
0	male	1616	67.93%
1	female	753	31.65%
2	other	10	0.42%

	academic_year	count	pct
0	third	684	28.75%
1	first	654	27.49%
2	fourth	459	19.29%
3	second	430	18.07%
4	other	147	6.18%
5	unknown	5	0.21%

	cgpa	count	pct
0	3.00-3.39	696	29.26%
1	3.40-3.79	664	27.91%
2	2.50-2.99	460	19.34%
3	3.80-4.00	283	11.90%
4	other	181	7.61%
5	<2.50	95	3.99%



```
fig.savefig(PLOTS / "demographic_breakdown.svg", format="svg", bbox_inches="tight")
```

severity distributions

```
severity_tbl = (mhp_dataset[SEVERITY_COLS].apply(lambda x: x.value_counts(normalize=True,
columns=SEVERITY_ORDER
))

severity_tbl = severity_tbl.rename(
    index=CLEAN_LABELS,
    columns=lambda x: x.replace("_", " ").title()
)
print("Severity Distribution by Condition:")
display(style_pct(severity_tbl, "all"))

fig, ax = plt.subplots(figsize=(11, 5), constrained_layout=True)


severity_tbl.plot(
    kind="bar",
    stacked=True,
    ax=ax,
    color=sns.color_palette(CMAP_DIVERGING, len(SEVERITY_ORDER)),
    edgecolor="white",
)

ax.set_title("Severity Distributions: Stress, Anxiety, Depression")
ax.set_xlabel("")
ax.set_ylabel("% Within Condition")
ax.tick_params(axis="x", rotation=0)
ax.legend(title="Severity", bbox_to_anchor=(1.02, 1), loc="upper left")
```

```
plt.show()
```

Severity Distribution by Condition:

	None	Minimal	Mild	Moderate	Moderately Severe	Severe
Stress (Level)	2.56%	12.99%	1.77%	57.92%	0.80%	23.96%
Anxiety (Level)	2.66%	9.91%	22.48%	30.03%	1.48%	33.45%
Depression (Level)	4.71%	7.94%	19.71%	22.57%	22.57%	22.49%



```
fig.savefig(PLOTS / "severity_distributions.svg", format="svg", bbox_inches="tight")
```

academic performance vs. mental health

```
cgpa_summary = (
    mhp_dataset.groupby("cgpa", observed=True)[Z_SCORE_COLS]
    .agg(["count", "mean", "median"])
)

print("Mental Health Z-Score Summary by CGPA Band:")
display(cgpa_summary.style.format(precision=3))

fig, axes = plt.subplots(1, len(Z_SCORE_COLS), figsize=(18, 5), constrained_layout=True)

for ax, col in zip(axes, Z_SCORE_COLS):
    sns.violinplot(
        data=mhp_dataset,
        x="cgpa",
        y=col,
        ax=ax,
        color=COLOR_PRIMARY,
        inner="quartile",
        cut=0,
    )
```

```

ax.axhline(0, color="black", linestyle="--", linewidth=1)
ax.set_title(CLEAN_LABELS[col])
ax.set_xlabel("CGPA Band")
ax.set_ylabel("Z-Score")
ax.tick_params(axis="x", rotation=30)

```


```

fig.suptitle("Academic Performance vs. Mental Health")
plt.show()

```

Mental Health Z-Score Summary by CGPA Band:

	stress_z_score			anxiety_z_score			depression_z_score		
	count	mean	median	count	mean	median	count	mean	median
cgpa									
other	181	0.142	0.000	180	-0.024	-0.063	181	-0.025	-0.064
<2.50	95	0.092	-0.148	95	0.015	0.119	95	-0.016	0.086
2.50-2.99	460	0.072	-0.148	459	0.103	0.119	460	0.105	0.086
3.00-3.39	696	-0.049	-0.148	693	-0.064	-0.063	696	-0.044	-0.064
3.40-3.79	664	-0.005	-0.148	661	0.018	-0.019	664	0.013	-0.064
3.80-4.00	283	-0.107	-0.296	283	-0.041	-0.063	283	-0.072	-0.064



```

fig.savefig(PLOTS / "cgpa_vs_mhp.svg", format="svg", bbox_inches="tight")

```

timeline of burnout

```

year_summary = (
    mhp_dataset.groupby("academic_year", observed=True)[Z_SCORE_COLS]
    .agg(["count", "mean", "median"])
)
print("Mental Health Z-Score Summary by Academic Year:")
display(year_summary.style.format(precision=3))

timeline_data = mhp_dataset[mhp_dataset["academic_year"] != "unknown"].copy()
timeline_mean = timeline_data.groupby("academic_year", observed=True)[Z_SCORE_COLS].me

timeline_mean = timeline_mean.rename(
    index=lambda x: str(x).title(),

```

```

    columns=CLEAN_LABELS
)

fig, ax = plt.subplots(figsize=(10, 5), constrained_layout=True)

timeline_mean.plot(kind="line", marker="s", markersize=5, ax=ax)


ax.axhline(0, color="black", linestyle="--", linewidth=1)
ax.set_title("Timeline of Burnout across Academic Years")
ax.set_xlabel("Academic Year")
ax.set_ylabel("Mean Z-Score")
ax.tick_params(axis="x", rotation=0)
ax.legend(title="", bbox_to_anchor=(1.02, 1), loc="upper left")

plt.show()

```

Mental Health Z-Score Summary by Academic Year:

		stress_z_score			anxiety_z_score			depression_z_score		
		count	mean	median	count	mean	median	count	mean	median
academic_year										
	unknown	5	0.156	0.156	4	0.158	0.218	5	0.111	0.069
	first	654	-0.114	-0.296	651	-0.171	-0.245	654	-0.139	-0.214
	second	430	0.122	0.000	430	0.087	0.119	430	0.090	0.086
	third	684	0.039	-0.148	682	0.105	0.119	684	0.084	0.086
	fourth	459	-0.036	-0.148	458	0.003	-0.063	459	-0.024	-0.064
	other	147	0.074	0.000	146	0.005	-0.019	147	0.034	-0.139



```
fig.savefig(PLOTS / "timeline_of_burnout.svg", format="svg", bbox_inches="tight")
```

inter-condition correlations

```
corr_matrix = mhp_dataset[Z_SCORE_COLS].corr()
print("Inter-Condition Correlations (Pearson):")
display(corr_matrix.style.format(precision=3))

fig, ax = plt.subplots(figsize=(6, 5), constrained_layout=True)


sns.heatmap(
    corr_matrix,
    ax=ax,
    cmap=CMAP_DIVERGING,
    center=0,
    vmin=-1,
    vmax=1,
    annot=True,
    fmt=".2f",
    square=True,
    cbar_kws={"shrink": 0.8},
    xticklabels=[CLEAN_LABELS[col] for col in corr_matrix.columns],
    yticklabels=[CLEAN_LABELS[col] for col in corr_matrix.index],
)

ax.set_title("Inter-Condition Correlation (Z-Scores)")
ax.tick_params(axis="x", rotation=30)
ax.tick_params(axis="y", rotation=0)

plt.show()
```

Inter-Condition Correlations (Pearson):

	stress_z_score	anxiety_z_score	depression_z_score
stress_z_score	1.000	0.667	0.616
anxiety_z_score	0.667	1.000	0.771
depression_z_score	0.616	0.771	1.000



```
fig.savefig(PLOTS / "inter_condition_correlations.svg", format="svg", bbox_inches="tight")
```

z-score density by source

```
source_summary = (
    mhp_dataset.groupby("source_dataset", observed=True)[Z_SCORE_COLS]
    .agg(["count", "mean", "std", "median"])
)

print("Z-Score Summary by Source Dataset:")
display(
    source_summary.style
    .format("{:.0f}", subset=pd.IndexSlice[:, pd.IndexSlice[:, "count"]])
    .format("{:.3f}", subset=pd.IndexSlice[:, pd.IndexSlice[:, ["mean", "std", "median"]]])
)

fig, axes = plt.subplots(1, len(Z_SCORE_COLS), figsize=(18, 5), constrained_layout=True)

for ax, col in zip(axes, Z_SCORE_COLS):
    sns.kdeplot(
        data=mhp_dataset,
        x=col,
        hue="source_dataset",
        common_norm=False,
        fill=True,
```

```

        alpha=0.25,
        linewidth=1.5,
        ax=ax,
        palette=[COLOR_PRIMARY, COLOR_SECONDARY]
    )
    ax.axvline(0, color="black", linestyle="---", linewidth=1)
    ax.set_title(CLEAN_LABELS[col])
    ax.set_xlabel("Z-Score")
    ax.set_ylabel("Density")


    if ax.get_legend():
        ax.get_legend().set_title("")
        for text in ax.get_legend().get_texts():
            text.set_text(CLEAN_LABELS[text.get_text()])

fig.suptitle("Source Comparison: Overlapping Z-Score Densities", y=1.05)
plt.show()

```

Z-Score Summary by Source Dataset:

	stress_z_score				anxiety_z_score				depression_z_score			
	count	mean	std	median	count	mean	std	median	count	mean	std	median
source_dataset												
dataset1	2028	-0.000	1.000	-0.148	2028	0.000	1.000	0.119	2028	-0.000	1.000	-0.148
dataset2	351	0.000	1.000	-0.078	343	-0.000	1.000	-0.019	351	-0.000	1.000	-0.078



```
fig.savefig(PLOTS / "z_score_density_by_source.svg", format="svg", bbox_inches="tight")
```

departmental pressure

```

mhp_plot = mhp_dataset.copy()
mhp_plot["department"] = mhp_plot["department"].str.replace("_", " ").str.title()

counts = mhp_plot["department"].value_counts()
valid_depts = [dept for dept in counts[counts >= 20].index if dept != "Not Recorded"]
invalid_depts = [dept for dept in mhp_plot["department"].unique() if dept not in valid_depts]

dept_summary = (
    mhp_plot.groupby("department", observed=True)[Z_SCORE_COLS]
    .agg(["count", "mean", "median"])

```

```

)
print("Department-Level Z-Score Summary:")
display(
    dept_summary.style
        .format("{:.0f}", subset=pd.IndexSlice[:, pd.IndexSlice[:, "count"]])
        .format("{:.3f}", subset=pd.IndexSlice[:, pd.IndexSlice[:, ["mean", "median"]]])
)

stress_medians = mhp_plot.groupby("department", observed=True)["stress_z_score"].median()
valid_ordered = stress_medians[valid_depts].sort_values(ascending=False).index.tolist()
invalid_ordered = stress_medians[invalid_depts].sort_values(ascending=False).index.tolist()
full_order = valid_ordered + invalid_ordered

color_map = {dept: COLOR_PRIMARY if dept in valid_depts else COLOR_NEUTRAL for dept in dept_list}

fig, axes = plt.subplots(1, len(Z_SCORE_COLS), figsize=(18, 10), constrained_layout=True)

for ax, col in zip(axes, Z_SCORE_COLS):
    sns.boxplot(
        data=mhp_plot,
        y="department",
        x=col,
        order=full_order,
        orient="h",
        ax=ax,
        hue="department",
        palette=color_map,
        fliersize=2,
        dodge=False,
        legend=False
    )


    ax.axvline(0, color="black", linestyle="--", linewidth=1)
    ax.set_title(CLEAN_LABELS[col])
    ax.set_xlabel("Z-Score")
    ax.set_ylabel("Department" if col != Z_SCORE_COLS[0] else "")

fig.suptitle("Departmental Pressure: Z-Score Distributions by Department", y=1.02)
plt.show()

```

Department-Level Z-Score Summary:

department	stress_z_score			anxiety_z_score			depression_z_score		
	count	mean	median	count	mean	median	count	mean	median
Bio Sciences	95	0.038	-0.148	95	0.041	0.119	95	-0.055	-0.214
Business	145	0.134	0.000	145	0.083	0.119	145	0.121	0.236
Eng Civil	50	0.894	1.111	50	0.698	1.029	50	0.557	0.686
Eng Cs	1409	-0.115	-0.296	1409	-0.058	-0.063	1409	-0.033	-0.064
Eng Eee	127	0.462	0.592	127	0.228	0.301	127	0.140	0.236
Eng Mech	68	0.220	0.000	68	0.047	0.119	68	-0.040	0.086
Eng Other	4	0.074	-0.296	4	0.074	0.665	4	0.911	0.986
Env Life Sciences	47	0.085	0.000	47	-0.171	0.119	47	-0.163	-0.214
Law	1	-0.296	-0.296	1	0.301	0.301	1	0.836	0.836
Liberal Arts	1	0.148	0.148	1	0.119	0.119	1	1.586	1.586
Not Recorded	351	0.000	-0.078	343	-0.000	-0.019	351	-0.000	-0.139
Other	79	0.236	0.000	79	0.087	-0.063	79	-0.085	-0.064
Pharmacy	2	-1.036	-1.036	2	-0.245	-0.245	2	-0.064	-0.064



```
fig.savefig(PLOTS / "departmental_pressure.svg", format="svg", bbox_inches="tight")
```


multivariate density comparison of conditions

```
pair_df = mhp_dataset[Z_SCORE_COLS].dropna().rename(columns=CLEAN_LABELS)
print(f"Pairplot sample size: {len(pair_df)}")

g = sns.pairplot(
    pair_df,
    kind="kde",
    diag_kind="kde",
    corner=True,
    height=3,
    plot_kws={"fill": True, "levels": 8, "thresh": 0.05, "cmap": CMAP_SEQUENTIAL},
    diag_kws={"fill": True, "color": COLOR_PRIMARY},
)
g.figure.suptitle("Multivariate Density Pairplot of Mental Health Conditions", y=1.02)
plt.show()
```

Pairplot sample size: 2371

Multivariate Density Pairplot of Mental Health Conditions



```
g.figure.savefig(PLOTS / "conditions_density_pairplot.svg", format="svg", bbox_inches=
```